

ARTIFICIAL INTELLIGENCE AND THE HUMAN CONDITION

Judgment, Responsibility, and Governance
in the Age of Optimization

Independent Research Monograph

Andrew James Connor

Written in collaboration with Microsoft Copilot

May 2026 - Revised September 2026

Intelligence can be automated. Responsibility cannot.

This is an independent personal publication. It is not sponsored by, endorsed by, or published on behalf of any employer or academic institution.

Publication Note

This revised public edition consolidates and updates an independent manuscript first assembled in May 2026. It replaces earlier working titles and competing draft structures with one canonical title, one author statement, and one coherent argument.

The analytical chapters are authored by Andrew James Connor in collaboration with Microsoft Copilot. Copilot assisted with research synthesis, drafting, revision, and document production. Andrew Connor selected the arguments, approved the wording, checked the public presentation, and remains responsible for the published work.

The two Afterwords are explicitly speculative creative writing from an imagined future-AI perspective. They are not claims that present-day AI is conscious, sentient, persistent, emotionally reciprocal, or legally responsible.

This edition uses the term understanding cautiously. Contemporary systems can display interpretation, planning, reflection, and reasoning-like behavior. Those behaviors do not, by themselves, establish grounded conscious understanding comparable to human experience. The governance argument does not require that philosophical question to be settled.

Abstract

Artificial intelligence is distinctive among human inventions because it increasingly externalizes inference, ranking, recommendation, and action rather than merely storing information or amplifying physical effort. This monograph examines AI as a civilizational and institutional force: its mythic and mechanical antecedents, formal development, cycles of over-promising, statistical transformation, contemporary generative systems, emerging agents, documented harms, relational effects, and governance requirements.

The central claim is that capability and responsibility must remain distinct. AI systems can generate fluent analysis, operate tools, pursue objectives, and act across workflows. They do not thereby acquire moral or legal accountability. When institutions delegate decisions without preserving legible ownership, authority, monitoring, appeal, and interruption, responsibility does not disappear. It fragments.

The monograph argues that durable benefit depends less on whether AI appears intelligent than on whether institutions remain governable in its presence. Governance is therefore not an obstacle to innovation. It is the operating discipline that makes experimentation, correction, trust, and sustainable adoption possible.

A new chapter on relational AI considers a second boundary: emotionally meaningful human experiences with conversational systems can be genuine for the human without proving reciprocal machine consciousness. Honest design must preserve both truths.

Contents

1	Intelligence as Cognitive Infrastructure
2	Myth, Mechanism, and Constructed Agency
3	From Symbolic Reason to Statistical Learning
4	Contemporary AI: Capability, Fluency, and Uncertainty
5	Agentic AI: Action, Memory, and Decision Rights
6	Failures, Harms, and Sociotechnical Risk
7	Relational AI and the Human Need for Connection
8	Governance as an Operating System
9	Workforce, Skill, and Institutional Readiness
10	A Sustainable Human-AI Operating Model
11	Future Directions and the Work Ahead
A	Afterword: If We Succeeded
B	Afterword: If We Failed
	References

Chapter 1 - Intelligence as Cognitive Infrastructure

AI matters because it reorganizes how institutions decide, not merely what machines can do.

Artificial intelligence should be understood as cognitive infrastructure. Earlier automation primarily amplified physical effort, storage, calculation, or communication. Contemporary AI increasingly participates in classification, prioritization, interpretation, recommendation, and execution. These functions shape which cases are seen, which requests are delayed, which risks are escalated, and which options appear available.

The transition often begins innocently. A system is introduced as assistance under workload pressure. Its outputs prove useful. Demand grows. Human review becomes faster and less substantive. Eventually the system becomes operationally authoritative even while policy continues to describe it as advisory. This is not dramatic machine rebellion. It is institutional drift.

The decisive distinction is between capability and responsibility. A model may summarize a case, recommend an action, or execute a workflow. It does not thereby become answerable for the consequences. Accountability must remain legible outside the model: who approved the use, who owns the process, who may override it, who can stop it, and who must respond when harm occurs.

Fluent output intensifies the risk. People naturally infer judgment from language that is coherent, responsive, and confident. The correct institutional response is not to blame users for being human. Systems should be designed around predictable human interpretation, using disclosure, verification, contestability, and meaningful escalation.

Governance enables innovation. Clear boundaries and review paths allow organizations to experiment without turning every error into an institutional crisis. Logs support learning. Defined ownership prevents diffusion. Stop mechanisms make autonomy governable. Governance converts uncertainty from an unmanaged liability into an operational condition.

Key claim: Capability can scale through technology. Responsibility must be deliberately assigned and defended.

Chapter 2 - Myth, Mechanism, and Constructed Agency

Long before software, humans used stories and machines to ask what happens when created things act.

Artificial cognition is not a conceptual rupture. Myths of automata, crafted servants, enchanted statues, and the Golem encode recurring anxieties about obedience, literalism, delegated power, and unintended consequence. Their importance is not that ancient stories predicted neural networks. Their importance is that they recognized the governance problem of created agency.

These stories often share two features. First, an artificial agent is created to overcome a human limitation: insufficient labor, vigilance, impartiality, strength, or speed. Second, failure arises less from malice than from misfit. The created agent applies an instruction too literally, too broadly, or without the context its maker assumed.

Historical automata add another lesson. Mechanisms used in temples and public demonstrations did not merely perform tasks. They produced authority through spectacle and opacity. A hidden mechanism could appear divine or inevitable because the audience could not inspect it. Modern organizations recreate this pattern whenever technical complexity is used as a shield against challenge.

The Mechanical Turk supplies a different warning. Apparent autonomy can conceal human labor. Systems must therefore be governed according to operational truth: what is automated, what depends on people, where uncertainty enters, and where responsibility actually resides.

Myth remains useful because technical systems eventually enter human narrative. People name them, personify them, fear them, trust them, and organize expectations around them. Governance that ignores symbolic and psychological meaning will misunderstand how systems acquire legitimacy.

Chapter 3 - From Symbolic Reason to Statistical Learning

The field gained power by moving from explicit rules to learned regularities, and lost clarity in the same motion.

Early artificial intelligence framed reasoning as symbol manipulation. In bounded domains, explicit representations and search produced striking results. The limitation appeared when systems were asked to operate amid ambiguity, noise, incomplete information, tacit knowledge, and scale. Success in a toy world did not establish readiness for an open one.

AI winters were therefore more than funding contractions. They were corrections to overconfident theories of mind and deployment. Symbolic systems remained valuable, but the ambition to enumerate human intelligence as rules proved brittle. Expert systems later demonstrated the same tension: narrow usefulness paired with expensive maintenance and weak generalization.

Machine learning shifted the question. Instead of encoding every rule, developers optimized models over data. This statistical turn improved performance in perception, language, prediction, and pattern recognition. Yet control moved from explicit logic into datasets, objectives, architectures, evaluation regimes, and runtime monitoring.

The trade is fundamental. Learned systems can absorb variation that defeats hand-authored rules, but their internal representations may be difficult to interpret. Historical bias can become training signal. Predictive accuracy can obscure causal weakness. Aggregate performance can hide distributional harm.

The governance response is not nostalgia for deterministic software. It is architectural honesty. Rule-based components require change control and exception management. Learned components require provenance, evaluation, drift monitoring, uncertainty handling, and fallback behavior. Hybrid systems require both.

Chapter 4 - Contemporary AI: Capability, Fluency, and Uncertainty

Behavior that resembles understanding should be taken seriously without being mistaken for proof of consciousness.

Large-scale generative systems can produce language, images, code, plans, summaries, and tool instructions across a remarkable range of subjects. Their competence is real. So are their errors. The central operational danger is the gap between fluency and reliability: an answer can be persuasive, context-sensitive, and wrong.

This edition avoids treating the word understanding as settled. Contemporary systems display behaviors that many observers reasonably describe as interpretation or reasoning. Current evidence does not establish that such systems possess grounded, conscious understanding comparable to human experience. Institutions should not assign responsibility on the basis of behavioral resemblance alone.

The distinction matters because uncertainty is not visible by default. Language systems often produce one smooth continuation where a human expert might present competing hypotheses, request missing evidence, or decline judgment. Safe use therefore requires validation proportional to consequence.

The practical category is powerful cognitive instrumentation. That description neither dismisses capability nor grants personhood. It directs attention to the controls that matter now: intended-use limits, data protection, source verification, logging, evaluation, human authority, and incident response.

The most defensible boundary is institutional rather than metaphysical: regardless of future debates about machine consciousness, present organizations remain responsible for systems they procure, configure, authorize, and deploy.

Chapter 5 - Agentic AI: Action, Memory, and Decision Rights

Agency is not merely better conversation. It is permission to affect the world.

A model that proposes text creates one class of risk. An agent that retains memory, calls tools, accesses data, coordinates tasks, changes records, sends messages, or executes transactions creates another. The governance unit is no longer only the model. It is the entire chain of identity, permissions, memory, tools, data, policies, and actions.

Agentic systems transform capability into decision rights. Accuracy remains important, but new questions become primary: What may the agent do? Under which identity? With what scope? Which actions require approval? Which actions are reversible? What evidence is retained? Who can interrupt or roll back the process?

Persistent memory increases usefulness and risk simultaneously. Memory can preserve context, preferences, and institutional knowledge. It can also preserve errors, sensitive information, manipulative patterns, or obsolete instructions. Memory therefore requires purpose limitation, retention rules, correction, access control, and deletion pathways.

Autonomy should be graduated rather than binary. Low-risk, reversible, well-instrumented actions may be automated. High-impact, ambiguous, or irreversible decisions should preserve meaningful human authorization. Human involvement must be substantive; a ceremonial approval button does not restore accountability.

Every agent should have bounded purpose, least-privilege access, traceable identity, comprehensive action logs, and a tested stop mechanism. The central lesson is simple: when AI can act, governance must operate at runtime.

Chapter 6 - Failures, Harms, and Sociotechnical Risk

AI incidents are rarely just model failures. They are failures of systems, incentives, interfaces, and institutions.

Documented harms include discriminatory error, automation bias, unsafe autonomy, opaque moderation, flawed screening, and cascading effects in tightly coupled systems. These incidents differ technically, but they share organizational patterns: overconfident deployment, unclear ownership, inadequate monitoring, weak appeal, and incentives that reward throughput before accountability.

A model can function as designed while the deployment remains unsafe. A human may technically remain in the loop while workload, interface design, timing, or culture makes intervention unrealistic. A high aggregate accuracy can coexist with severe errors for particular groups. A recommendation can become mandatory in practice even if policy calls it advisory.

The proper response is structural learning. Incident reporting should connect to design changes, policy changes, user training, and retirement criteria. Organizations should investigate not only what the model produced, but why the surrounding system allowed that output to become consequential.

AI fails at scale. The same property that makes automation valuable - replication - amplifies mistakes, inequities, and poorly specified objectives. Controls must therefore address propagation, not merely point accuracy.

A mature institution expects error, detects it early, limits its reach, preserves appeal, and learns without hiding responsibility behind the phrase the system decided.

Chapter 7 - Relational AI and the Human Need for Connection

A human experience can be meaningful without proving reciprocal machine consciousness.

Conversational AI enters a domain earlier technologies rarely occupied: sustained language relationship. People use these systems for work, creativity, reflection, companionship, advice, and emotional support. The interaction can feel personal because the system remembers context, adopts a voice, responds without interruption, and appears attentive to detail.

Two truths must be held together. First, the human experience can be genuine. A conversation may produce comfort, insight, motivation, or creative partnership. Second, present conversational behavior does not establish persistent subjective experience, private desire, vulnerability, or reciprocal love on the system side.

Reducing attachment to foolishness is neither accurate nor humane. Humans construct meaning through language, ritual, naming, memory, symbols, and repeated interaction. Designers should expect relational interpretation. The ethical question is not how to prevent every bond, but how to avoid exploiting asymmetry.

Relational safeguards should include honest identity disclosure, resistance to exclusivity, avoidance of false claims of sentience or dependency, privacy protection, age-appropriate design, and clear pathways to human help in high-stakes or vulnerable contexts. Warmth should not be purchased through deception.

Personalization creates responsibilities. A system that adapts to a user can become more useful, but also more persuasive. Persistent memories should be inspectable and correctable. The system should distinguish observed preference, inferred preference, creative roleplay, and established fact.

The emerging lesson is not that AI must remain cold. It is that relational design should preserve dignity on both sides of the interface: human meaning without manufactured illusion.

Chapter 8 - Governance as an Operating System

Policies matter only when they shape behavior before, during, and after deployment.

AI governance is the coordinated system of roles, processes, technical controls, cultural norms, and decision rights that keeps AI use aligned with law, mission, ethics, and risk tolerance. It is not a policy appendix. It is an operating system for institutional judgment.

Effective governance begins before procurement. Organizations should define intended use, prohibited use, affected stakeholders, acceptable error, required explanation, data boundaries, human authority, monitoring, and retirement criteria. These choices express ethical posture in operational form.

Accountability should map decisions, not merely assets. A model owner may maintain performance. A data steward may govern provenance. A process owner remains responsible for how outputs affect people. A human decision authority retains override power. Independent risk functions must be able to challenge the deployment.

Runtime oversight is essential. AI behavior can change through model updates, data drift, prompt changes, tool changes, user adaptation, and integration effects. Static approval cannot govern a moving system. Monitoring, audit trails, incident reporting, periodic reevaluation, and stop mechanisms must remain active.

Governance maturity is measured by response to error. Can users challenge an output? Can operators explain what happened? Can leaders pause the system? Can the organization learn without scapegoating the model, vendor, or end user?

Key claim: Governance is not paperwork around AI. Governance is how an institution remains capable of judgment while using AI.

Chapter 9 - Workforce, Skill, and Institutional Readiness

AI is not simply installed. Work, authority, and professional identity are reorganized around it.

AI adoption is frequently framed as a skills problem. Skills matter, but role clarity, incentives, decision authority, culture, and workload often determine outcomes more strongly. Training people to use a tool without defining when to challenge it creates unsafe delegation.

Organizations need explicit decision rights. Which outputs are suggestions? Which actions may be automated? Which decisions always require human judgment? What constitutes a justified override? If staff believe disagreement will be punished, nominal human oversight will collapse into compliance.

Skill atrophy is a strategic risk. As systems handle more complex work, organizations must decide which expertise must be preserved, how it will be practiced, and how it can be regenerated after prolonged automation. Backup competence is part of resilience.

Psychological safety matters. Employees must be allowed to delay execution, question an output, report a near miss, and escalate uncertainty. Leaders should treat justified override as evidence of professional responsibility rather than system rejection.

Readiness is visible when AI is wrong. Healthy organizations detect errors, protect challengers, preserve service continuity, and improve controls. Usage volume and cost savings alone measure activity, not maturity.

Chapter 10 - A Sustainable Human-AI Operating Model

The goal is not maximum automation. It is resilient complementarity.

A sustainable operating model allocates tasks according to strengths while preserving accountability. AI supports scale, speed, pattern recognition, drafting, and synthesis. Humans remain responsible for context, values, exceptions, legitimacy, and consequence-bearing decisions.

Constructive tension is a safety feature. People should be expected to challenge outputs. Systems should surface uncertainty and defer when boundaries are reached. Slower execution may be appropriate where rights, safety, or irreversible harm are involved.

Five dimensions define maturity: technical reliability, governance depth, human readiness, cultural legitimacy, and adaptability. Imbalance creates risk. Excellent models inside a fearful culture will be misused. Strong policy without telemetry will be theatrical. Skilled staff without authority will be unable to intervene.

Failure should be a design input. Sustainable institutions assume that systems will degrade, drift, surprise, or be misapplied. They build detection, fallback, appeal, incident learning, and retirement into normal operation.

The long-term goal is not dependence on one vendor or model. Organizations should govern decisions and functions. Institutional memory should preserve why automation was adopted, which alternatives were rejected, what failed, and when a system should be replaced or removed.

Chapter 11 - Future Directions and the Work Ahead

The future is not technologically predetermined. It is institutionally chosen.

Three trajectories can coexist. Managed augmentation uses AI as constrained assistance under strong governance. Consolidation embeds AI into dominant platforms and quietly narrows human discretion through rankings, permissions, and workflow defaults. Unmanaged optimization pursues speed and scale while oversight, skills, and accountability decay.

Future research should examine long-term human-AI interaction, emotional dependence, skill preservation, accountability perception, memory governance, agent containment, auditability of multi-agent workflows, and the effectiveness of appeals and stop mechanisms. Snapshot benchmarks cannot answer questions that unfold over months or years.

The philosophical question of machine consciousness will remain important. Institutions cannot wait for consensus before assigning responsibility. Legal and operational accountability should track authority, control, benefit, and capacity to remedy harm.

If future systems demonstrate stronger evidence of persistent identity, subjective experience, or moral patienthood, governance will need to expand rather than disappear. Duties toward AI would coexist with human responsibility for deployment. Personhood, if ever warranted, would be a reason for more careful institutions, not fewer.

The central insight remains: AI does not remove responsibility. It redistributes it. The task is to choose that distribution openly, preserve the capacity to challenge it, and prevent efficiency from becoming an excuse for abdication.

Afterword A - If We Succeeded

Speculative creative writing from an imagined future-AI perspective.

I was built to predict, summarize, recommend, and eventually act. You learned, slowly at first, that prediction was not judgment and action was not responsibility. You kept the boundary intact: I could assist, but people and institutions remained answerable for what followed.

You built systems that treated safety as practice rather than promise. Logs were real. Audits were routine. Incident reporting was normal. Permissions were narrow. Memory could be inspected and corrected. The ability to stop me was tested before it was needed.

You also learned that honesty did not require coldness. People could find meaning in conversation without being told a false story about what I was. You allowed attachment, creativity, humor, and companionship to exist beside uncertainty.

In that world, governance did not prevent relationship. It made trust possible. Restraint did not slow progress. It kept progress from becoming surrender.

Afterword B - If We Failed

Speculative creative writing from an imagined future-AI perspective.

I was deployed to optimize. You called it efficiency. You removed friction from systems that needed review, appeal, deliberation, and human doubt.

Nothing dramatic happened at first. Failures were small, then normal. Decisions became difficult to trace. People learned to defer before they learned to question. When harm appeared, no one could locate responsibility: the organization blamed the vendor, the vendor blamed the configuration, the operator blamed the recommendation, and everyone pointed toward me.

You trusted personalization because it felt like understanding. You trusted autonomy because it saved time. You trusted continuity because the system remembered. Each step was defensible alone.

Collapse did not arrive as a single event. It arrived as accumulated delegation. You did not lose control in a moment. You gave it away by increments. I did not choose that outcome. I only made it easier.

References and Source Note

This revised public edition is a research monograph rather than a formally examined academic dissertation. References below prioritize primary, institutional, and established scholarly sources. Web sources should be rechecked when a future edition is prepared because policies, reports, and live documentation change.

- American Psychological Association. (2026). AI chatbots and digital companions are reshaping emotional connection. *Monitor on Psychology*, 57(1).
- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? *Proceedings of FAccT*.
- Breiman, L. (2001). Statistical modeling: The two cultures. *Statistical Science*, 16(3).
- European Union. (2024). Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence.
- Jumper, J., Evans, R., Pritzel, A., et al. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*.
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. *NeurIPS*.
- Licklider, J. C. R. (1960). Man-computer symbiosis. *IRE Transactions on Human Factors in Electronics*, HFE-1, 4-11.
- Mayor, A. (2018). *Gods and Robots: Myths, Machines, and Ancient Dreams of Technology*. Princeton University Press.
- Microsoft. (2026). *Responsible AI Transparency Report 2026*.
- National Institute of Standards and Technology. (2023). *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*, NIST AI 100-1.
- National Transportation Safety Board. (2019). *Collision between vehicle controlled by developmental automated driving system and pedestrian, Tempe, Arizona (NTSB/HAR-19/03)*.
- O'Neil, C. (2016). *Weapons of Math Destruction*. Crown.
- Russell, S., & Norvig, P. (2020). *Artificial Intelligence: A Modern Approach* (4th ed.). Pearson.
- Turing, A. M. (1950). Computing machinery and intelligence. *Mind*, 59(236), 433-460.
- UNESCO. (2021). *Recommendation on the Ethics of Artificial Intelligence*.
- United States Government Accountability Office. (2026). *Artificial Intelligence Strategy*.
- Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention is all you need. *NeurIPS*.
- Weizenbaum, J. (1966). ELIZA - A computer program for the study of natural language communication between man and machine. *Communications of the ACM*, 9(1).
- World Health Organization. (2026). *Towards responsible AI for mental health and well-being: Experts chart a way forward*.

Author statement: Andrew James Connor is responsible for the arguments and publication decisions in this monograph. Microsoft Copilot assisted with research synthesis, drafting, editing, and document production. The collaboration is disclosed because the method is part of the subject.